

3: AGENT ARCHITECTURE

AI6125: Intelligent Agents

Assoc Prof Zhang Jie

Agent Architecture

- An *agent architecture* is a *software design* for an agent
- We have already seen a top-level decomposition (abstract), into:
 perception – state – decision – action
- An agent architecture defines (concrete):
 - Key data structures;
 - Operations on data structures;
 - Control flow between operations

History of Agent Architectures

- Originally (1956-1985), pretty much all agents designed within AI were ***symbolic reasoning agents***
- Its purest expression proposes that agents use *explicit logical reasoning* in order to decide what to do
- Problems with symbolic reasoning led to a reaction against this — the so-called ***reactive agents*** movement, 1985–present
- From 1990-present, a number of alternatives proposed: ***hybrid architectures***, which attempt to combine the best of reasoning and reactive architectures

no long-term planning.

Symbolic Reasoning Agents

- The classical approach to building agents is to view them as a particular type of knowledge-based system, and bring all the associated (discredited?!) methodologies of such systems to bear
- This paradigm is known as *symbolic AI*
- We define a ***deliberative*** agent or agent architecture to be one that:
 - contains an explicitly represented, symbolic model of the world
 - makes decisions (for example about what actions to perform) via symbolic reasoning

Symbolic Reasoning Agents

- If we aim to build an agent in this way, there are **two key problems to be solved:**

1. *The transduction problem:*
that of translating the real world into an accurate, adequate symbolic description, in time for that description to be useful...vision, speech understanding, learning
2. *The representation/reasoning problem:*
that of how to symbolically represent information about complex real-world entities and processes, and how to get agents to reason with this information in time for the results to be useful...knowledge representation, automated reasoning, automatic planning

Symbolic Reasoning Agents

- Most researchers accept that neither problem is anywhere near solved
- Underlying problem lies with the complexity of symbol manipulation algorithms in general: many (most) search-based symbol manipulation algorithms of interest are *highly intractable*
- Because of these problems, some researchers have looked to alternative techniques for building agents; we look at these later

Deductive Reasoning Agents

- How can an agent decide what to do using theorem proving?
- Basic idea is to use logic to encode a theory stating the *best* action to perform in any given situation
- Let:
 - ρ be this theory (typically a set of rules) *facts*
 - Δ be a logical database that describes the current state of the world
 - Ac be the set of actions the agent can perform
 - $\Delta \vdash_{\rho} \phi$ mean that ϕ can be proved from Δ using ρ

Deductive Reasoning Agents

/ try to find an action explicitly prescribed */*

for each $a \in Ac$ do

 if $\Delta \vdash_{\rho} Do(a)$ then
 return a

 end-if

end-for

/ try to find an action not excluded */*

for each $a \in Ac$ do

 if $\Delta \not\vdash_{\rho} \neg Do(a)$ then
 return a


 end-if

end-for

return *null* */* no action found */*

Deductive Reasoning Agents

- An example: The Vacuum World
- Goal is for the robot to clear up all dirt



Deductive Reasoning Agents

- Use 3 *domain predicates* to solve problem:

$In(x, y)$ agent is at (x, y)

$Dirt(x, y)$ there is dirt at (x, y)

$Facing(d)$ the agent is facing direction d

- Possible actions:

$Ac = \{turn, forward, suck\}$

P.S. *turn* means “turn right”

Deductive Reasoning Agents

- Rules ρ for determining what to do:

$In(0, 0) \wedge Facing(north) \wedge \neg Dirt(0, 0) \longrightarrow Do(forward)$

$In(0, 1) \wedge Facing(north) \wedge \neg Dirt(0, 1) \longrightarrow Do(forward)$

$In(0, 2) \wedge Facing(north) \wedge \neg Dirt(0, 2) \longrightarrow Do(turn)$

$In(0, 2) \wedge Facing(east) \longrightarrow Do(forward)$

- ...and so on!
- Using these rules (+ other obvious ones), starting at (0, 0) the robot will clear up dirt

Deductive Reasoning Agents

■ Problems:

- ❑ How to convert video camera input to $Dirt(0, 1)$?
- ❑ decision making assumes a static environment: *calculative* rationality
- ❑ decision making using first-order logic is *undecidable*!

■ Typical solutions:

- ❑ weaken the logic
- ❑ use symbolic, non-logical representations
- ❑ shift the emphasis of reasoning from *run time* to *design time*

↑
NP-hard.

AGENT0 and PLACA

- Much of the interest in agents from the AI community has arisen from Shoham's notion of *agent oriented programming* (AOP)
- AOP: 'a new programming paradigm, based on a societal view of computation'
- The key idea that informs AOP is that of directly programming agents in terms of intentional notions like belief, commitment, and intention

AGENT0

- AGENT0 is implemented as an extension to LISP
- Each agent in AGENT0 has 4 components:
 - a set of capabilities (things the agent can do)
 - a set of initial beliefs
 - a set of initial commitments (things the agent will do)
 - a set of commitment rules 类似于therapy.
- The key component, which determines how the agent acts, is the commitment rule set

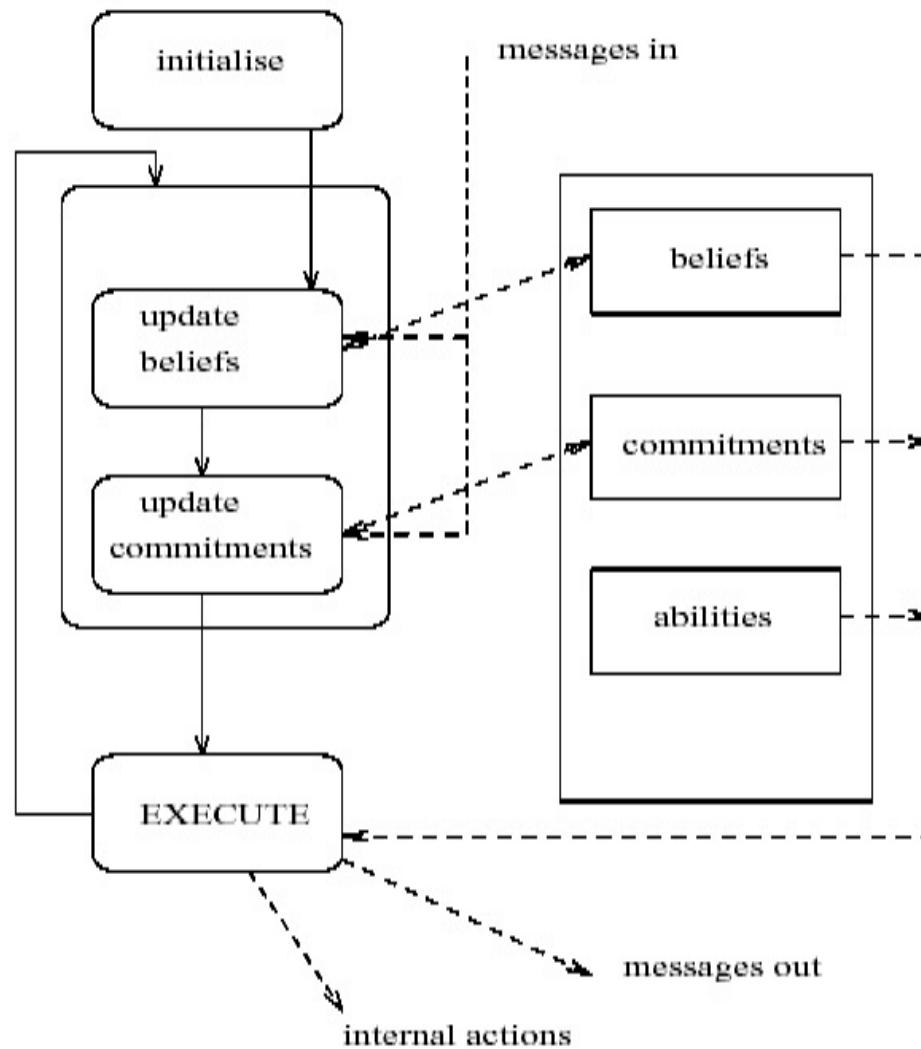
AGENT0

- Each commitment rule contains
 - *a message condition*
 - *a mental condition*
 - an action
- On each ‘agent cycle’ ...
 - The message condition is matched against the messages the agent has received
 - The mental condition is matched against the beliefs of the agent
 - If the rule fires, then the agent becomes committed to the action (the action gets added to the agent’s commitment set)

AGENT0

- Actions may be
 - private:
an internally executed computation, or
 - communicative:
sending messages
- Messages are constrained to be one of three types:
 - “requests” to commit to action
 - “unrequests” to refrain from actions
 - “informs” which pass on information

AGENT0



AGENT0

- A commitment rule:

```
COMMIT(  
  ( agent, REQUEST, DO(time, action)  
  ), ;;; msg condition 1  
  ( B,  
    [now, Friend agent] AND  
    CAN(self, action) AND  
    NOT [time, CMT(self, anyaction)]  
  ), ;;; mental condition 2  
  self,  
  DO(time, action) 3  
)
```

AGENT0

- This rule may be paraphrased as follows:
if I receive a message from *agent* which requests me to do *action* at *time*, and I believe that:
 - *agent* is currently a friend
 - I can do the action
 - At *time*, I am not committed to doing any other action
then commit to doing *action* at *time*

AGENT0 and PLACA

- AGENT0 provides support for multiple agents to cooperate and communicate, and provides basic provision for debugging...
- ...it is, however, a *prototype*, that was designed to illustrate some principles, rather than be a production language
- A more refined implementation was developed by Thomas, for her 1993 doctoral thesis
- Her Planning Communicating Agents (PLACA) language was intended to address one severe drawback to AGENT0: the inability of agents to plan, and communicate requests for action via high-level goals
- Agents in PLACA are programmed in much the same way as in AGENT0, in terms of *mental change* rules

AGENT0 and PLACA

- An example mental change rule:

```
((self ?agent REQUEST (?t (xeroxed ?x)))  
  (AND (CAN-ACHIEVE (?t xeroxed ?x))  
        (NOT (BEL (*now* shelving)))  
        (NOT (BEL (*now* (vip ?agent)))))  
  ((ADOPT (INTEND (5pm (xeroxed ?x))))))  
  ((?agent self INFORM  
    (*now* (INTEND (5pm (xeroxed ?x)))))))
```

- Paraphrased:

if someone asks you to xerox something, and you can, and you don't believe that they're a VIP, or that you're supposed to be shelving books, then

- adopt the intention to xerox it by 5pm, and
- inform them of your newly adopted intention

Practical Reasoning

- Practical reasoning is reasoning directed towards actions — the process of figuring out what to do:
 - “Practical reasoning is a matter of weighing conflicting considerations for and against competing options, where the relevant considerations are provided by what the agent desires/values/cares about and what the agent believes.” (Bratman)
- Practical reasoning is distinguished from *theoretical reasoning* – theoretical reasoning is directed towards beliefs

Practical Reasoning

- Human practical reasoning consists of two activities:
 - 1 □ *deliberation*
deciding *what* state of affairs we want to achieve
 - 2 □ *means-ends reasoning*
deciding *how* to achieve these states of affairs
- The outputs of deliberation are *intentions*

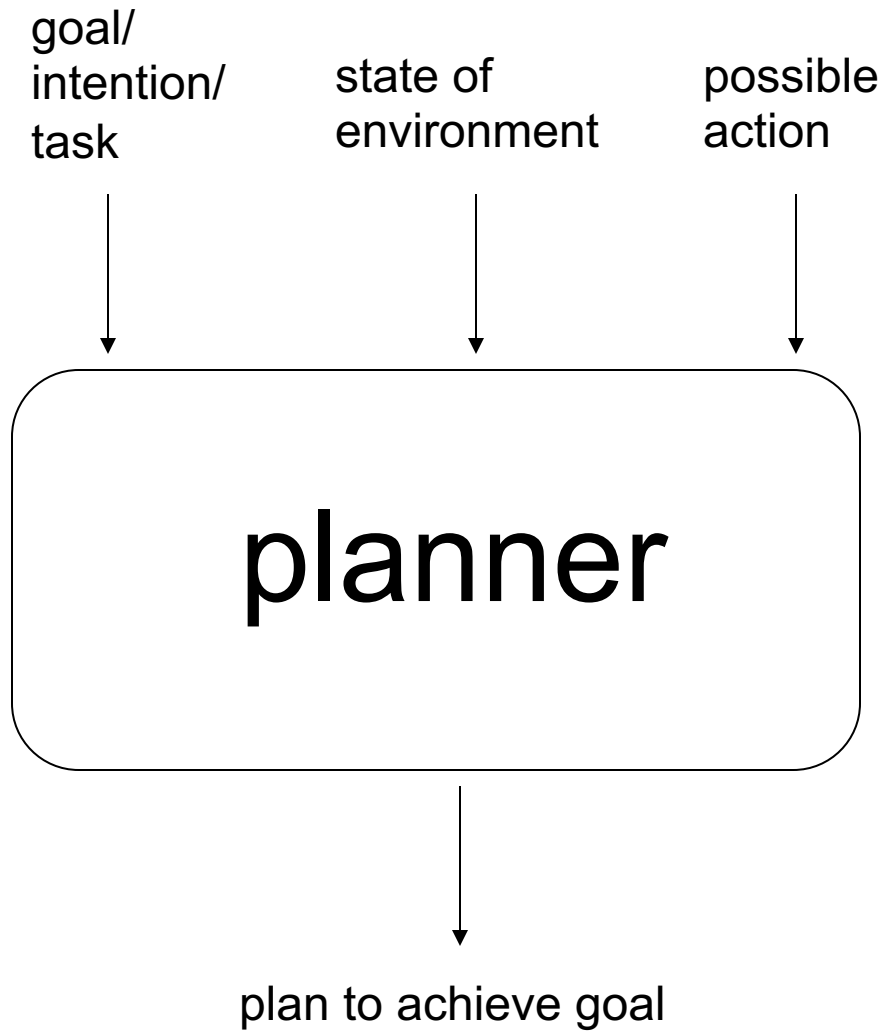
Intentions in Practical Reasoning

- Notice that intentions are much stronger than mere desires:

“My desire to play basketball this afternoon is merely a potential influencer of my conduct this afternoon. It must vie with my other relevant desires [. . .] before it is settled what I will do. In contrast, once I intend to play basketball this afternoon, the matter is settled: I normally need not continue to weigh the pros and cons. When the afternoon arrives, I will normally just proceed to execute my intentions.” (Bratman, 1990)

What is Means-End Reasoning?


- Basic idea is to give an agent:
 - representation of goal/intention to achieve
 - representation of actions it can perform
 - representation of the environmentand have it generate a *plan* to achieve the goal
- Essentially, this is
automatic programming



Planning

- Question: How do we *represent*. . .
 - goal to be achieved
 - state of environment
 - actions available to agent
 - plan itself

The Blocks World



- We'll illustrate the techniques with reference to the *blocks world* (like previous module)
- Contains a robot arm, 3 blocks (A, B, and C) of equal size, and a table-top

The Blocks World Ontology

- To represent this environment, need an *ontology*

On(x , y)

obj x on top of obj y

OnTable(x)

obj x is on the table

Clear(x)

nothing is on top of obj x

Holding(x)

arm is holding x

The Blocks World


- Here is a representation of the blocks world described above:

Clear(A)

On(A, B)

OnTable(B)

OnTable(C)




- Use the *closed world assumption*: anything not stated is assumed to be *false*

The Blocks World

- A *goal* is represented as a set of formulae
- Here is a goal:

$$OnTable(A) \wedge OnTable(B) \wedge OnTable(C)$$




The Blocks World

- *Actions* are represented using a technique that was developed in the STRIPS planner
- Each action has:
 - a *name*
which may have arguments
 - a *pre-condition list*
list of facts which must be true for action to be executed
 - a *delete list*
list of facts that are no longer true after action is performed
 - an *add list*
list of facts made true by executing the action

Each of these may contain *variables*

The Blocks World Operators




■ Example:

The stack action occurs when the robot arm places the object x it is holding on top of object y .

	$Stack(x, y)$
pre	$Clear(y) \wedge Holding(x)$
del	$Clear(y) \wedge Holding(x)$
add	$ArmEmpty \wedge On(x, y)$

A Plan



■ What is a plan?

A sequence (list) of actions, with variables replaced by constants.

Implementing Practical Reasoning Agents

- A first pass at an implementation of a practical reasoning agent:

Agent Control Loop Version 1

```
1. while true
2.     observe the world;
3.     update internal world model;
4.     deliberate about what intention to achieve next;
5.     use means-ends reasoning to get a plan for the intention;
6.     execute the plan
7. end while
```

- (We will not be concerned with stages (2) or (3))

Implementing Practical Reasoning Agents

- Let's make the algorithm more formal:

```
Agent Control Loop Version 2
1.   $B := B_0$ ; /* initial beliefs */
2.  while true do
3.      get next percept  $\rho$ ;
4.       $B := brf(B, \rho)$ ;
5.       $I := deliberate(B)$ ;
6.       $\pi := plan(B, I)$ ;
7.      execute( $\pi$ )
8.  end while
```

Deliberation

- How does an agent deliberate?
 - begin by trying to understand what the *options* available to you are
 - *choose between them*, and *commit* to some
- Chosen options are then intentions

Deliberation

- The *deliberate* function can be decomposed into two distinct functional components:

- 1 □ *option generation*

in which the agent generates a set of possible alternatives;

Represent option generation via a function, *options*, which takes the agent's current beliefs and current intentions, and from them determines a set of options (= *desires*)

- 2 □ *filtering*

in which the agent chooses between competing alternatives, and commits to achieving them.

In order to select between competing options, an agent uses a *filter* function.

Deliberation

BDI agents

Agent Control Loop Version 3

```
1.
2.    $B := B_0;$    belief
3.    $I := I_0;$    intention
4.   while true do
5.       get next percept  $\rho$ ;
6.        $B := brf(B, \rho);$ 
7.        $D := options(B, I);$ 
8.        $I := filter(B, D, I);$ 
9.        $\pi := plan(B, I);$ 
10.      execute( $\pi$ )
11.  end while
```

Reactive Architectures


- There are many unsolved (some would say insoluble) problems associated with symbolic AI
- These problems have led some researchers to question the viability of the whole paradigm, and to the development of *reactive* architectures
- Although united by a belief that the assumptions underpinning mainstream AI are in some sense wrong, reactive agent researchers use many different techniques
- We start by reviewing the work of one of the most vocal critics of mainstream AI: Rodney Brooks

Rodney A. Brooks

- M.I.T professor
- Member of M.I.T.'s Artificial Intelligence Lab
- Developed the Subsumption Architecture for robot control in 1986
- His goal was to develop artificial, complete creatures capable of inhabiting our world, not a simplified world


The Subsumption Architecture

- The Subsumption Architecture is:
 - A layering methodology for robot control systems
 - A parallel and distributed method for connecting sensors and actuators in robots



The Subsumption Architecture

- Each layer is made up of connected, simple processors: Augmented Finite State Machines



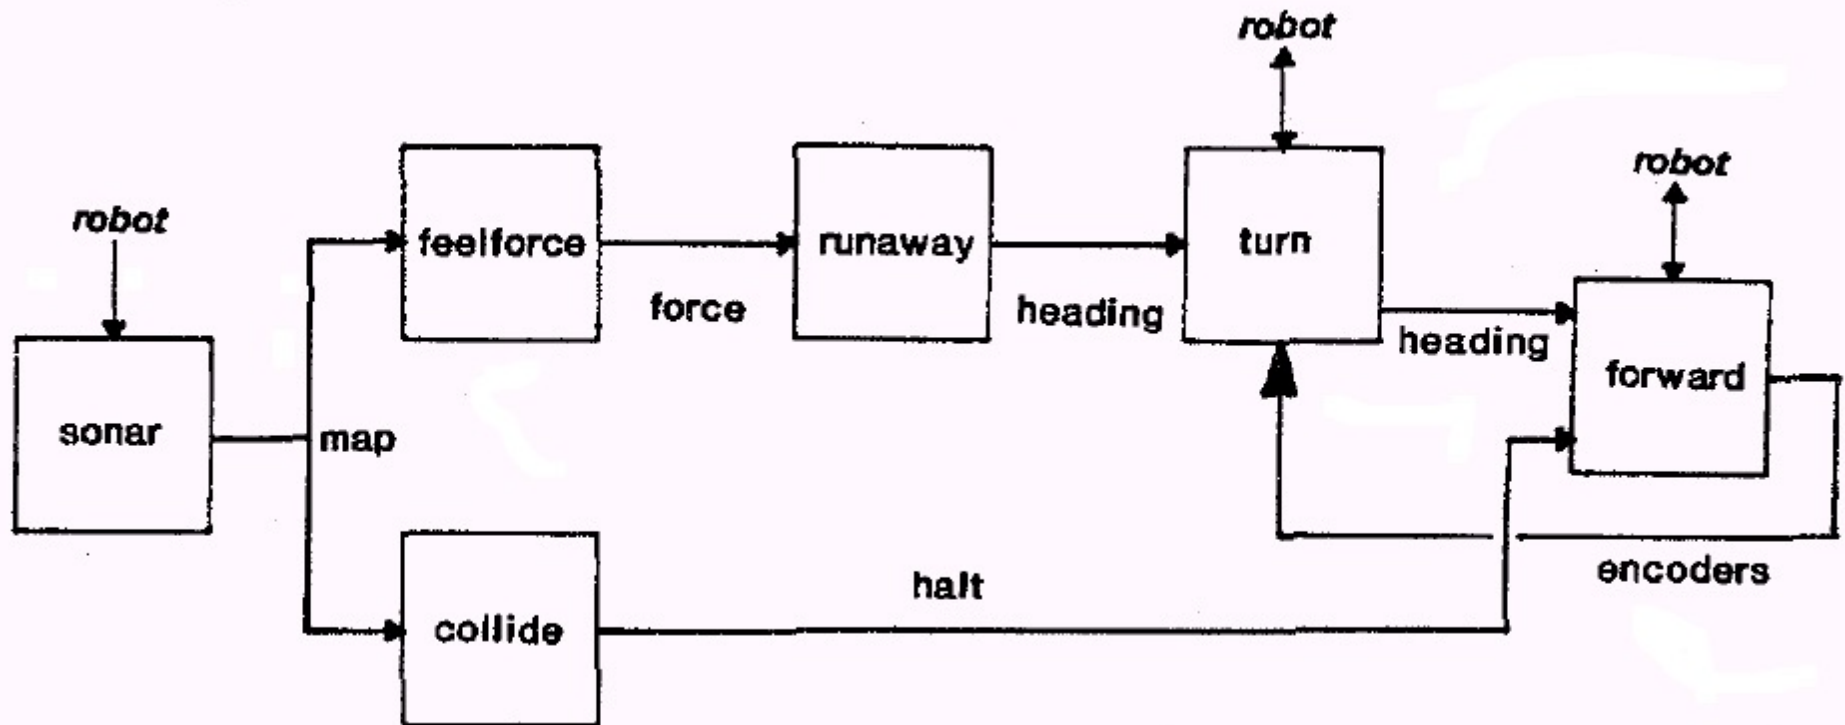
The Subsumption Architecture

- The most important aspect of these FSMs
 - Outputs are simple functions of inputs and local variables
 - Inputs can be suppressed and outputs can be inhibited
 - This function allows higher levels to subsume the function of lower levels
 - Lower, therefore, still function as they would without the higher levels

An Example: Allen

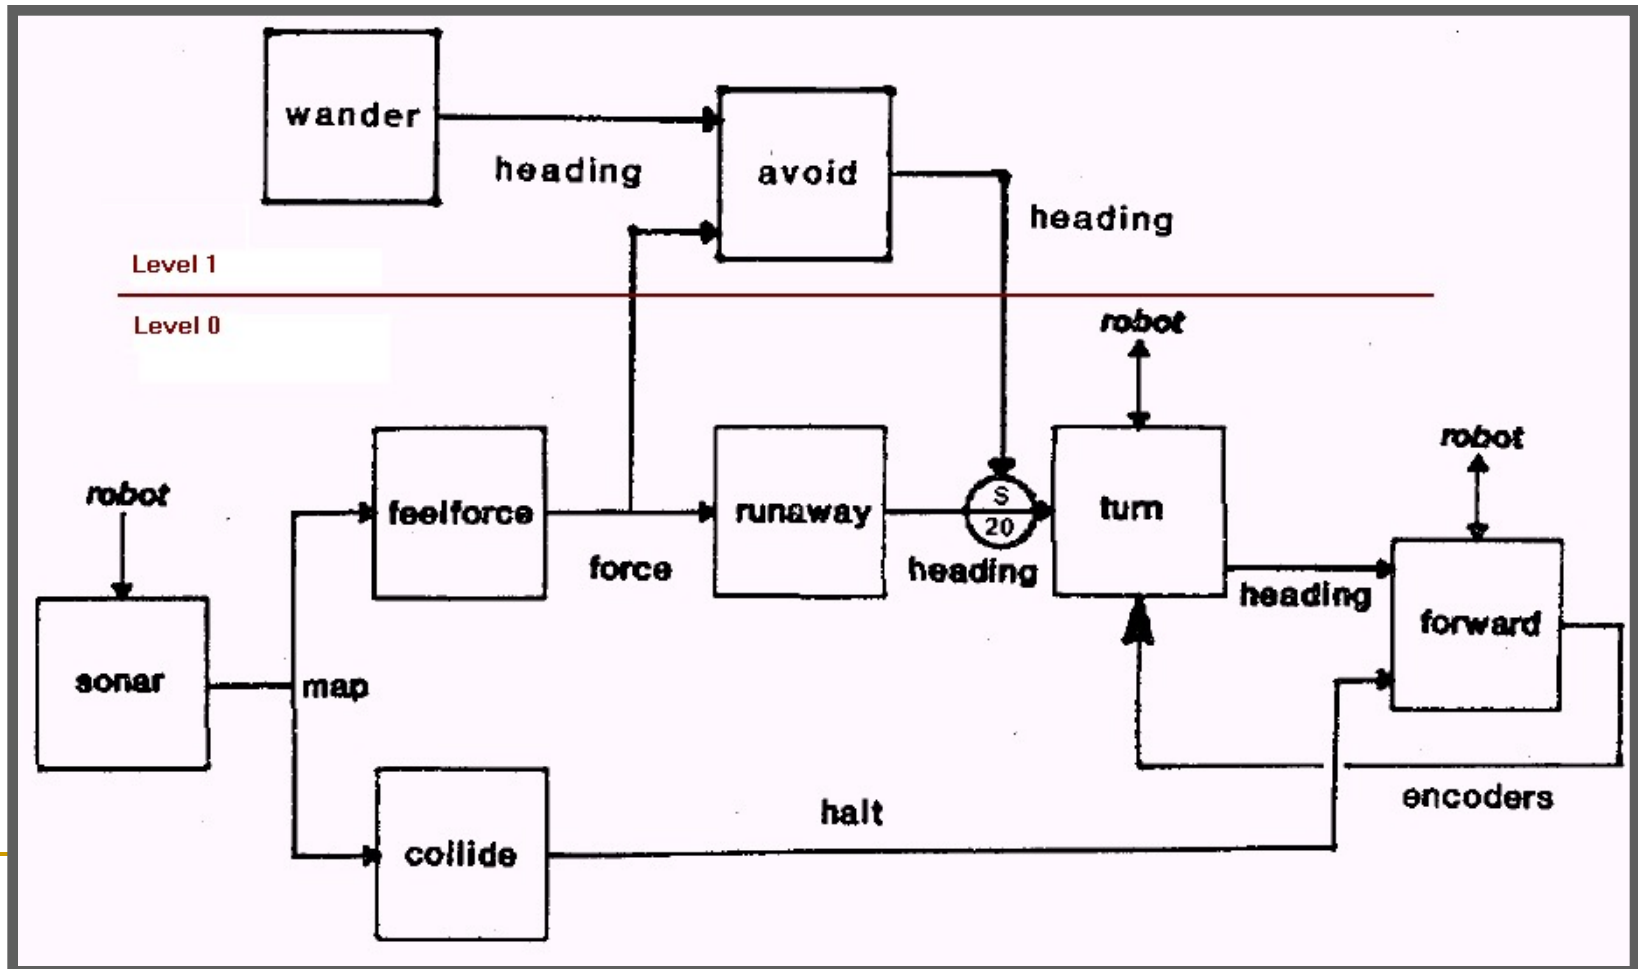


- Brooks' first Subsumption robot
- Level 0: Runs away if approached, avoids objects



An Example: Allen

- Levels 1 and 0: Adds wandering



- Levels 2, 1, and 0: Adds hallway following



Programming Characteristics of Subsumption

- No internal model of the real world because:
 - No free communication
 - No shared memory
- So, use real world as the model
 - “The world really is a rather good model of itself”¹
 - Very accurate
 - Never out of date
 - No computation needed to keep model up to date
- Real world used for sub-system communication
 - Instead of direct communication, sub-systems just sense the real world

Brooks – behavior languages

- Brooks has put forward three theses:
 1. Intelligent behavior can be generated *without* explicit representations of the kind that symbolic AI proposes
 2. Intelligent behavior can be generated *without* explicit abstract reasoning of the kind that symbolic AI proposes
 3. Intelligence is an *emergent* property of certain complex systems

Brooks – behavior languages

- He identifies two key ideas that have informed his research:
 1. Situatedness and embodiment: ‘Real’ intelligence is situated in the world, not in disembodied systems such as theorem provers or expert systems
 2. Intelligence and emergence: ‘Intelligent’ behavior arises as a result of an agent’s interaction with its environment. Also, intelligence is ‘in the eye of the beholder’ ; it is not an innate, isolated property

Brooks – behavior languages

- To illustrate his ideas, Brooks built some robots based on his *subsumption architecture*
- A subsumption architecture is a hierarchy of task-accomplishing *behaviors*
- Each behavior is a rather simple rule-like structure
- Each behavior ‘competes’ with others to exercise control over the agent
- Lower layers represent more primitive kinds of behavior (such as avoiding obstacles), and have precedence over layers further up the hierarchy
- The resulting systems are, in terms of the amount of computation they do, *extremely* simple
- Some of the robots do tasks that would be impressive if they were accomplished by symbolic AI systems

Steels' Mars Explorer

- Steels' Mars explorer system, using the subsumption architecture, achieves near-optimal cooperative performance in simulated 'rock gathering on Mars' domain:

The objective is to explore a distant planet, and in particular, to collect sample of a precious rock. The location of the samples is not known in advance, but it is known that they tend to be clustered.

Steels' Mars Explorer Rules

- For individual (non-cooperative) agents, the lowest-level behavior, (and hence the behavior with the highest “priority”) is obstacle avoidance:
if detect an obstacle then change direction (1)
- Any samples carried by agents are dropped back at the mother-ship:
*if carrying samples and at the base
then drop samples* (2)
- Agents carrying samples will return to the mother-ship:
*if carrying samples and not at the base
then travel up gradient* (3)

Steels' Mars Explorer Rules

- Agents will collect samples they find:
if detect a sample then pick sample up (4)
- An agent with “nothing better to do” will explore randomly:
if true then move randomly (5)

Advantages of Reactive Agents

- Simplicity
- Economy
- Computational tractability
- Robustness against failure
- Elegance

Limitations of Reactive Agents

- Agents without environment models must have sufficient information available from local environment
- If decisions are based on *local* environment, how does it take into account *non-local* information (i.e., it has a “short-term” view)
- Difficult to make reactive agents that learn
- Since behavior emerges from component interactions plus environment, it is hard to see how to *engineer* specific agents (no principled methodology exists)
- It is hard to engineer agents with large numbers of behaviors (dynamics of interactions become too complex to understand)

Hybrid Architectures

- Many researchers have argued that neither a completely deliberative nor completely reactive approach is suitable for building agents
- They have suggested using *hybrid* systems, which attempt to marry classical and alternative approaches
- An obvious approach is to build an agent out of two (or more) subsystems:
 - a *deliberative* one, containing a symbolic world model, which develops plans and makes decisions in the way proposed by symbolic AI
 - a *reactive* one, which is capable of reacting to events without complex reasoning

Hybrid Architectures

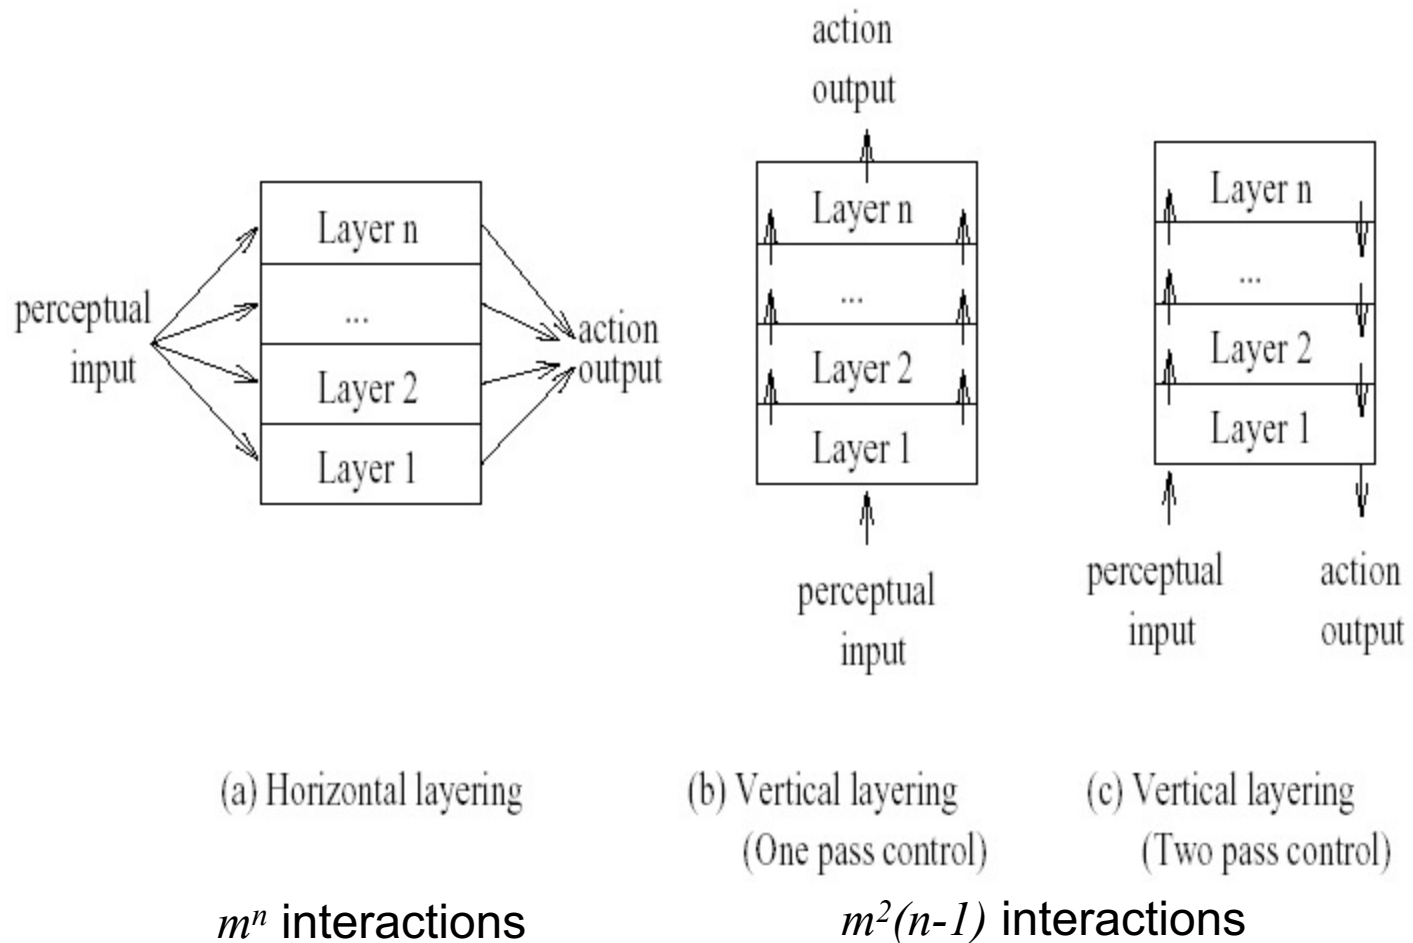
- Often, the reactive component is given some kind of precedence over the deliberative one
- This kind of structuring leads naturally to the idea of a *layered* architecture, of which TOURINGMACHINES and INTERRAP are examples
- In such an architecture, an agent's control subsystems are arranged into a hierarchy, with higher layers dealing with information at increasing levels of abstraction

Hybrid Architectures

- A key problem in such architectures is what kind of control framework to embed the agent's subsystems in, to manage the interactions between the various layers
- Horizontal layering
Layers are each directly connected to the sensory input and action output.
In effect, each layer itself acts like an agent, producing suggestions as to what action to perform.
- Vertical layering
Sensory input and action output are each dealt with by at most one layer each

Hybrid Architectures

m possible actions suggested by each layer, n layers



Introduces bottleneck
in central control system

Not fault tolerant to
layer failure

Summary

- Symbolic Reasoning Agents
 - Deductive Reasoning Agents
 - Agent Oriented Programming (AOP)
 - Practical Reasoning Agents
 - Problems with symbolic reasoning agents
- Reactive Agents
 - Limitations of reactive agents
- Hybrid Architectures